

AI Transparency Report

APRIL 2025



Table of Contents

Introduction.....3

I. Description of AI card4

 EU AI Act Risk Rating4

II. Categorization5

 a) OneTrust Copilot5

 b) Data Clustering – Document Clustering6

 b) Data Clustering – Document Classification7

 c) Auto-complete Responses – Record Analysis8

 c) Auto-complete Responses – Document Scanner.....9



DISCLAIMER:

No part of this document may be reproduced in any form without the written permission of OneTrust.

The contents of this document may be revised by OneTrust in its sole discretion, without notice, due to continued progress in the methodology of the Certification, any changes in applicable laws, regulations or related guidance, or for any other reason. OneTrust shall have no liability for any error or damage of any kind resulting from the use of this document, its contents or the information provided therewith.

The contents of this document, any materials and other information conveyed during this Privacy Automation Certification are for informational purposes only and do not constitute legal advice (and should not be relied upon as such).

Introduction

At OneTrust, trust is at the center of our values and that means we are committed to the responsible development and use of artificial intelligence (AI). We are dedicated to ensuring that AI deployed in our application (which we will refer to as AI features) not only meets regulatory requirements, but also aligns with industry best practices and OneTrust's AI Principles. OneTrust's AI Principles guide how our teams leverage AI within our solutions:

1. **Transparency by design:** We prioritize openness and explainability through in-product notices where appropriate and updates to our AI Transparency Report.
2. **Privacy-centric data protection:** Transparent data collection, accountable use, and responsible management of data are the key pillars of protecting data in AI features.
3. **Algorithmic accountability:** We require a clear system of accountability for managing AI risks.
4. **Fairness & inclusivity:** We strive to mitigate the risks of unwanted biases and discriminatory practices in our AI features. We believe AI should empower everyone and contribute to the overall prosperity, inclusivity, and growth for all.
5. **Data Stewardship:** Our AI features are designed with industry-recognized information security safeguards.
6. **Safety and reliability assurance:** Our AI features are designed to perform safely and reliably.

OneTrust released our first AI Transparency Report in 2024, and we continue to provide updates as we innovate, learn, and refine our practices. This AI Transparency Report provides information about AI features we have deployed through 'AI feature cards' that outline important information about each AI feature, such as its purpose, methodology, and risk rating. The AI feature cards are organized into categories based on their functionality and data used in development and function.

For more information on our approach to responsible AI see our [Trust page](#).



I. Description of AI card

Name of the AI feature, solution, and product: the name of the AI feature is at the top of the AI feature card. The solution(s) and the product(s) the AI feature is associated with are listed in the first row. In some cases, the AI features are applied platform wide.

Purpose: describes the use case and function of the AI feature.

Methodology: provides the name of the AI model used for the AI feature, details on the operation of the AI feature, and information on any training data used in development.

Data Processed: outlines the data used in execution to deliver the AI feature. Categories include public and proprietary sources and also customer data. Customer data (known also as customer content) means any data or information submitted by or on behalf of a OneTrust customer to the OneTrust platform and applications.

Customer data is never used to train the AI models.

Controls: details the options regarding customer control to enable or disable the AI feature.

- **Default:** indicates whether the AI feature is on or off by default (i.e. as the pre-selected option).
- **Option to Disable:** indicates if the customer administrator can turn the feature on or off.

Controls for enabling/ disabling AI features are found under Settings > Support > AI Feature Management.

EU AI Act Risk Rating

As part of our commitment to ethical AI development, OneTrust does not develop 'Prohibited AI systems' as defined under the EU AI Act and is not planning to offer any high-risk AI features or systems as part of the services.

Below are the descriptions of each AI risk category based on the risk qualification criteria outlined in the EU AI Act:

High-risk: AI features that pose a high probability of harm to health and safety or an adverse impact on the fundamental rights of individuals. Article 6 of the EU AI Act describes the thresholds that lead to an AI system being "high-risk." Either such a system meets the safety component criteria in Article 6(1) or falls into a category referred to in Annex III (Specific List of High-Risk AI systems).

Limited (transparency) risk: AI features that use General Purpose AI models to create original content as well as AI features that interact directly with natural persons. Examples include autonomous chatbots; technology that summarizes long-form content, autonomously creates software code, generates digital images from natural language, and produces articles or creative stories based on given prompts; music composition tools that create original scores and songs; and video generation tools that produce animated sequences or edit videos based on text descriptions.

Minimal risk: AI features that do not fall into High or Limited/ Transparency AI risk categories. Examples include spam filters, inventory management systems, and business process improvement tools.

II. Categorization

We organize our AI features into categories based on their functionality and similar data use.

a) OneTrust Copilot

This category includes AI features that deliver knowledge and answer customer questions by responding to input with data from trusted OneTrust sources such as product documentation and regulatory guidance. These AI features use natural language interfaces to deliver customized information, oftentimes with links to supporting material included for deeper reference and validation.

These AI features do not use any customer data apart from the prompts typed in by the customer to drive feature execution.

AI feature: OneTrust Copilot (AI Chat Assistant)	
Solution(s) and Product(s)	OneTrust Platform
Purpose	Conversational chatbot that answers customers' custom questions about global laws and regulations and how to use the OneTrust platform. Responses are generated from my.OneTrust product documentation and proprietary OneTrust DataGuidance research data. By delivering customized answers to specific customer questions, this AI feature saves customers time on research and finding information across sources.
Methodology	Which model is used? Azure OpenAI text-embedding-3-small (third-party), and Azure OpenAI GPT-4 variants (third-party) How does it work? A retrieval-augmented generation (RAG) technique converts DataGuidance and my.OneTrust data into vector representations using Azure OpenAI text-embedding-3-small and then compares the data to find relevant responses to the user's question. Training Data None.
Data Processed	my.OneTrust product documentation and DataGuidance research data is used. Customer input is processed exclusively to execute the model.
Controls	Default: On Option to disable: Yes
EU AI Act Risk Rating	Limited risk

b) Data Clustering – Document Clustering

These AI features evaluate and organize large sets of data by identifying similarities and patterns. They automatically group, classify, or categorize items — such as documents or records — to streamline data analysis, discovery, and compliance. This category of AI features evaluates customer data without storing or training on it.

AI feature: Document Clustering	
Solution(s) and Product(s)	Data Use Governance: Data Discovery & Classification
Purpose	Categorizes documents into groups based on perceived characteristics, thereby helping users identify personal and sensitive data, relevant topics, emerging trends, and outliers within their datasets.
Methodology	Which model is used? FastText 0.9.2 How does it work? The model processes unstructured data — such as files, articles, or reports — in a variety of formats. It analyzes the semantic data of each document and groups them into clusters based on similarity in meaning. Training Data Does not train on customer data.
Data Processed	Customer data in the form of documents identified during the Data Discovery scanning process.
Controls	Default: Off Option to disable: Yes
EU AI Act Risk Rating	Minimal risk

b) Data Clustering – Document Classification

AI feature: Document Classification	
Solution(s) and Product(s)	Data Use Governance: Data Discovery & Classification
Purpose	Categorizes documents into pre-defined types that are common in businesses. As an example, it can determine whether a document is a resume, patent, or contract. Document organization and retrieval is streamlined through this rapid classification process.
Methodology	Which model is used? FCustom support vector machine (SVM) based model (Onetrust proprietary), Word2Vec or GloVe for embeddings (third-party), and Tesseract for OCR (third-party) How does it work? The document type is predicted using supervised machine learning (ML) on text extracted from the file. This process is accomplished by the following steps: the extracted text is first converted to vector representation using text embedding techniques (such as word2vec or GloVe), which are pretrained or fine-tuned; next, based on the text embedding, a SVM classification algorithm is used to decide the document class. Where there are images, Tesseract is used for OCR. Training Data Does not train on customer data.
Data Processed	Customer data in the form of documents found during the Data Discovery scanning process.
Controls	Default: Off Option to disable: Yes
EU AI Act Risk Rating	Minimal risk

c) Auto-complete Responses – Record Analysis

This category includes AI features that automatically fill in responses using customer-provided documents or other information supplied by the customer. These tools reduce customers’ manual work by suggesting or pre-populating answers based on relevant context. These features temporarily evaluate customer data without storing or training on it.

AI feature: Inventory Record Analysis	
Solution(s) and Product(s)	Privacy Automation: Assessment Automation Third-Party Management: Third Party Risk Management
Purpose	Provides tailored recommendations for assessment questionnaires by analyzing responses from previously completed assessments.
Methodology	Which model is used? Azure OpenAI GPT-4 variants (third-party) How does it work? Retrieval augmented generation (RAG) is used for this feature. The system first evaluates the list of unanswered questions for an assessment. For each unanswered question, it identifies semantically similar questions from past assessments completed in the same client tenant and incorporates the responses into the final prompt for the large language model (LLM). The system utilizes the LLM's reasoning capabilities to recommend answers for approval by the user. Training Data None.
Data Processed	Customer data in the form of previously completed assessments.
Controls	Default: Off Option to disable: Yes
EU AI Act Risk Rating	Minimal risk

c) Auto-complete Responses – Document Scanner

AI feature: Document Scanner	
Solution(s) and Product(s)	Third-Party Management: Third Party Risk Management Privacy Automation: Assessment Automation
Purpose	Automatically populates assessment questionnaires with suggested responses using information in documents uploaded and scanned by the customer such as a project summary, contract, or SOC2 report. Saves the customer time and reduces customers' manual effort by auto-filling assessments using insights from documents added by the customer.
Methodology	Which model is used? Azure OpenAI ada-embedding-002 (third-party), and Azure OpenAI GPT-4 variants (third-party) How does it work? Retrieval augmented generation (RAG) is used for this functionality. Uploaded documents are turned into vector representations using Azure OpenAI ada-embedding-002 and compared to find the ones relevant to the user's assessment question. Training Data None.
Data Processed	Customer data in the form of documents that the customer chooses to upload during an assessment.
Controls	Default: Off Option to disable: Yes
EU AI Act Risk Rating	Minimal risk

The OneTrust logo, featuring the word "onetrust" in a white, lowercase, sans-serif font. The background is a dark green gradient with a large, lighter green circular shape on the right side.

OneTrust's mission is to enable the responsible use of data and AI. Our platform simplifies the collection of data with consent and preferences, automates the governance of data with integrated risk management across privacy, security, IT/tech, third-party, and AI risk, and activates the responsible use of data by applying and enforcing data policies across the entire data estate and lifecycle. OneTrust supports seamless collaboration between data teams and risk teams to drive rapid and trusted innovation. Recognized as a market pioneer and leader, OneTrust boasts over 300 patents and serves more than 14,000 customers globally, ranging from industry giants to small businesses. For more information, visit www.onetrust.com.